

Used Workbook 4 as the base. Just swapped out the base for this. Just swapped out the datasets.

```
In [1]: import warnings
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn import datasets
from sklearn.preprocessing import StandardScaler, scale
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split, GridSearchCV, cr
oss_val_score, cross_val_predict
from sklearn.metrics import confusion_matrix, precision_score, recall_
score, f1_score, roc_curve, roc_auc_score
from sklearn.metrics import precision_recall_curve, classification_rep
ort
```

```
In [2]: matches = pd.read_csv('Champion_Info1.csv', sep=",")
```

```
In [3]: matches.head()
```

Out[3]:

	gameID	platformID	seasonID	Winner	C1	C1S1	C1S2	C2	C2S1	C2S2	...	C7S2
0	3591397591	NA1	13	1	201	4	14	102	12	4	...	7
1	3653371287	NA1	13	0	120	11	6	13	12	4	...	4
2	3650198359	NA1	13	1	102	11	4	122	4	6	...	4
3	3645932094	NA1	13	0	18	4	7	99	3	4	...	4
4	3645761570	NA1	13	0	102	11	4	82	4	12	...	4

5 rows × 34 columns



```
In [4]: matches.shape
```

Out[4]: (965, 34)

```
In [5]: #Data Matrix
X = pd.DataFrame(matches, columns=[ "C1", "C1S1", "C1S2", "C2", "C2S1", "C2S2", "C3", "C4S1", "C4S2",
                                     "C4", "C4S1", "C4S2", "C5", "C5S1", "C5S2", "C6", "C6S1", "C6S2",
                                     "C7", "C7S1", "C7S2", "C8", "C8S1", "C7S2", "C9", "C9S1", "C9S2",
                                     "C10", "C10S1", "C10S2" ])

#Target Vector
y=matches[ 'Winner' ]

print(X.shape)
print(y.shape)
```

```
(965, 30)
(965,)
```

```
In [6]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
print(y_train.shape)
```

```
(772,)
```

```
In [7]: # Set the the range of K
neighbors = np.arange(1,60)

# Two arrays to store training and test accuracies
train_accuracy = np.empty(len(neighbors))
validation_accuracy = np.empty(len(neighbors))

for i,k in enumerate(neighbors):

    # Setup a knn classifier with k neighbors
    knn = KNeighborsClassifier(n_neighbors=k)

    # Fit the model
    knn.fit(X_train, y_train)

    # The "score" function returns the mean accuracy on the given train/test data and labels.
    # Note that "accuracy" may not be a good performance measure in a skewed data set
    # Thus, we need to do hyperparameter tuning by using better performance measures (e.g., f1 score, presision, recall)

    # Compute training accuracy
    train_accuracy[i] = knn.score(X_train, y_train)

    # Compute validation accuracy using cross-validation

    scores = cross_val_score(knn, X_train, y_train, scoring='accuracy', cv=5)


    validation_accuracy[i] = scores.mean()
```

```
In [8]: plt.style.use('ggplot')
fig = plt.figure(figsize=(10, 6))
plt.title('Varying K: Number of Nearest Neighbors')
plt.plot(neighbors, validation_accuracy, label='Validation Accuracy')
plt.plot(neighbors, train_accuracy, label='Training accuracy')
plt.legend()
plt.xlabel('K (Number of Nearest Neighbors)')
plt.ylabel('Accuracy')
plt.show()

# Find the value of "K" that gives max validation accuracy
j = 0
max_val_accuracy = validation_accuracy[j]
max_k = 1

for i in neighbors:
    if(validation_accuracy[j] > max_val_accuracy):
        max_val_accuracy = validation_accuracy[j]
        max_k = i
    j +=1

print("Optimal K: ", max_k)
```



Optimal K: 55

```
In [9]: %%time

warnings.filterwarnings('ignore')

# The param_grid tells Scikit-Learn to evaluate all combinations of the hyperparameter values
param_grid = {'n_neighbors': np.arange(1,50), 'p': [1, 2, 10, 50, 100, 500, 1000],
              'weights': ["uniform", "distance"]}

knn_clf = KNeighborsClassifier()

knn_cv = GridSearchCV(knn_clf, param_grid, scoring='f1', cv=5, verbose=3, n_jobs=-1)
knn_cv.fit(X_train, y_train)

params_optimal_knn = knn_cv.best_params_

print("Best Score: %f" % knn_cv.best_score_)
print("Optimal Hyperparameter Values: ", params_optimal_knn)
print("\n")
```

Fitting 5 folds for each of 686 candidates, totalling 3430 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n_jobs=-1)]: Done 24 tasks	elapsed: 35.2s
[Parallel(n_jobs=-1)]: Done 212 tasks	elapsed: 40.1s
[Parallel(n_jobs=-1)]: Done 852 tasks	elapsed: 56.5s
[Parallel(n_jobs=-1)]: Done 1748 tasks	elapsed: 1.3min
[Parallel(n_jobs=-1)]: Done 2900 tasks	elapsed: 1.9min
[Parallel(n_jobs=-1)]: Done 3430 out of 3430	elapsed: 2.1min finished

Best Score: 0.662806

Optimal Hyperparameter Values: {'n\_neighbors': 49, 'p': 1, 'weights': 'uniform'}

CPU times: user 3.93 s, sys: 342 ms, total: 4.27 s

Wall time: 2min 8s

```
In [10]: # With the Mahalanobis distance metric only the brute force algorithm
works
#knn = KNeighborsClassifier(weights=?, algorithm='brute', n_neighbors=
?, metric = "mahalanobis", metric_params= {'V': ?})

# Minkowski distance metric based optimal model selected via hyperpara
meter tuning.
# The Minkowski distance based model (i.e., knn_cv) is already trained
with the optimal hyperparameter values.
# We can use the optimal model (knn_cv) for prediction.
# Or we can use the optimal hyperparameter values to train a new model
, as follows.

knn = KNeighborsClassifier(**params_optimal_knn)

knn.fit(X_train, y_train)

y_train_predicted = knn.predict(X_train)
print(y_train_predicted)

train_accuracy_knn = np.mean(y_train_predicted == y_train)
print("\nTraining Accuracy: ", train_accuracy_knn)
```

```
[0 1 1 1 1 1 1 0 1 1 1 0 0 1 1 1 1 1 0 1 1 0 1 0 0 1 0 1 1 1 0 1 1 0
1 1 1
1 1 1 1 1 1 0 1 1 1 1 1 0 1 1 1 1 1 1 0 1 1 1 1 0 0 1 0 0 1 1 0 1 0
1 0 1
0 1 0 0 0 0 1 0 0 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 0 1
1 1 1
1 1 1 1 1 1 1 1 1 0 1 0 0 0 1 1 1 0 1 1 1 1 0 1 1 0 0 1 1 1 1 0 0 0
1 0 1
1 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
0 1 1
1 1 0 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 0 0 1 1 1 1 0 1 1 1 1 1 1 1 1 1
1 1 1
1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 0 0 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1
1 1 1 0 1 1 1 1 0 1 1 1 0 0 0 1 1 1 1 1 0 0 1 0 1 1 1 1 1 0 1 1 1 1
0 1 1
1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 0 1 1 0 1 1 1 1 1 0 1 1
1 1 1
1 1 1 1 1 1 1 1 1 1 1 1 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1
1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 0 1 1 0 1 1 1 1 1 1 1 1 1 0 1 1
0 1 1
1 0 1 0 1 0 1 0 1 1 1 1 0 1 1 1 1 1 1 1 1 0 1 1 0 1 1 1 1 1 1 0 1 1
0 1 1
1 1 1 0 1 1 1 1 1 1 1 1 1 0 1 1 1 1 0 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1
0 1 0
0 1 1 1 1 1 1 0 1 1 1 1 0 0 0 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 0 0 1 0 0
1 1 1
1 1 1 1 1 1 0 0 0 1 1 0 1 1 1 0 1 1 0 1 1 1 1 1 0 0 0 1 1 1 1 1 1 1
0 1 0
1 1 1 0 0 1 1 1 1 1 1 1 0 1 1 0 1 1 1 0 1 1 1 1 0 1 1 1 1 1 1 1 0 1
1 1 0
1 0 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 0 1 1 0 0 1 1 1 1 0
1 1 1
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 0 1 1 1 1 1
0 0 1
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 0 1 1 1 1 1 1 1 0 0 1 1 1 1 1 1
1 1 1
1 1 1 1 0 1 1 0 1 0 1 1 1 1 1 1 1 1 0 1 1 1 1 1 0 1 1 1 1 1 1 1 0
1 1 0
0 1 1 1 1 1 0 1 1 1 0 0 1 1 1 1 1 1 1 0 1 1 1 0 1 1 0 1 0 0 1 1]
```

Training Accuracy: 0.5764248704663213

```
In [11]: # Scoring Parameter for Classification:
# https://scikit-learn.org/stable/modules/model_evaluation.html#scoring-parameter
# Note: For a skewed data set "accuracy" might not be a good choice for scoring
scores = cross_val_score(knn, X_train, y_train, scoring='f1', cv=5)
print(scores)

print("F1 Score: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))

[0.66009852 0.70531401 0.63809524 0.65686275 0.65365854]
F1 Score: 0.66 (+/- 0.05)
```

```
In [12]: y_train_pred = cross_val_predict(knn, X_train, y_train, cv=5)

confusion_matrix(y_train, y_train_pred)
```

```
Out[12]: array([[ 84, 276],
               [ 71, 341]])
```

```
In [13]: precision = precision_score(y_train, y_train_pred)
print("Precision = %f" % precision)

recall = recall_score(y_train, y_train_pred)
print("Recall = %f" % recall)

f1 = f1_score(y_train, y_train_pred)
print("F1 Score = %f" % f1)

Precision = 0.552674
Recall = 0.827670
F1 Score = 0.662779
```



```
In [14]: #The accuracy of the model
train_accuracy_knn = knn.score(X_train, y_train)
print("\nTest Accuracy: ", train_accuracy_knn)

# No. of Correct Predictions
y_train_predicted = knn.predict(X_train)
print("\nNo. of correct predictions (Test): %d/%d" % (np.sum(y_train_p
redicted == y_train), len(y_train)))

# Confusion Matrix
print("\nConfusion Matrix (Train Data):\n", confusion_matrix(y_train,
y_train_predicted))

precision = precision_score(y_train, y_train_predicted)
print("Train Precision = %f" % precision)

recall = recall_score(y_train, y_train_predicted)
print("Train Recall = %f" % recall)

f1 = f1_score(y_train, y_train_predicted)
print("Train F1 Score = %f" % f1)
print(classification_report(y_train,y_train_predicted))
```

Test Accuracy: 0.5764248704663213

No. of correct predictions (Test): 445/772

Confusion Matrix (Train Data):

```
[[ 93 267]
 [ 60 352]]
```

Train Precision = 0.568659

Train Recall = 0.854369

Train F1 Score = 0.682832

	precision	recall	f1-score	support
0	0.61	0.26	0.36	360
1	0.57	0.85	0.68	412
accuracy			0.58	772
macro avg	0.59	0.56	0.52	772
weighted avg	0.59	0.58	0.53	772

```
In [15]: # Get the 2nd column of the matrix of predicted probabilities for each data point
# The 2nd column stores the probabilities of the positive class
y_scores = cross_val_predict(knn, X_train, y_train, method="predict_proba", cv=5)[: , 1]

fpr, tpr, thresholds = roc_curve(y_train, y_scores)

print("\nFPR FPR & TPR for Various Threshold Values:")

print("FPR: ", fpr)
print("TPR: ", tpr)
print("\nThresholds: ", thresholds)
```

FPR FPR & TPR for Various Threshold Values:


```
FPR: [0.          0.00277778 0.00833333 0.03333333 0.05833333 0.1055
5556
0.16666667 0.26666667 0.39722222 0.48611111 0.6          0.69166667
0.76666667 0.86666667 0.91666667 0.96111111 0.98611111 0.99166667
1.          1.          ]
TPR: [0.          0.00242718 0.00485437 0.01699029 0.04368932 0.0946
6019
0.14805825 0.25          0.34466019 0.48300971 0.59708738 0.72815534
0.8276699  0.90048544 0.94660194 0.9684466  0.98058252 0.99271845
0.99514563 1.          ]
```

```
Thresholds: [1.73469388 0.73469388 0.71428571 0.69387755 0.67346939
0.65306122
0.63265306 0.6122449  0.59183673 0.57142857 0.55102041 0.53061224
0.51020408 0.48979592 0.46938776 0.44897959 0.42857143 0.40816327
0.3877551  0.34693878]
```

```
In [16]: plt.style.use('ggplot')

fig = plt.figure(figsize=(10, 6))
def plot_roc_curve(fpr, tpr, label=None):
    plt.plot(fpr, tpr, color='darkorange', linewidth=8, label=label)
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.axis([0, 1, 0, 1])
    plt.title('ROC Curve (Train Data)')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')

plot_roc_curve(fpr, tpr)
plt.show()
```



```
In [17]: #Area under ROC curve
roc_auc_score(y_train,y_scores)
```

```
Out[17]: 0.5011023462783172
```

```

In [18]: plt.style.use('ggplot')

# Get the 2nd column of the matrix of predicted probabilities for each
data point
# The 2nd column stores the probabilities of the positive class
y_scores = cross_val_predict(knn, X_train, y_train, method="predict_proba", cv=3)[: , 1]

precisions, recalls, thresholds = precision_recall_curve(y_train, y_scores)

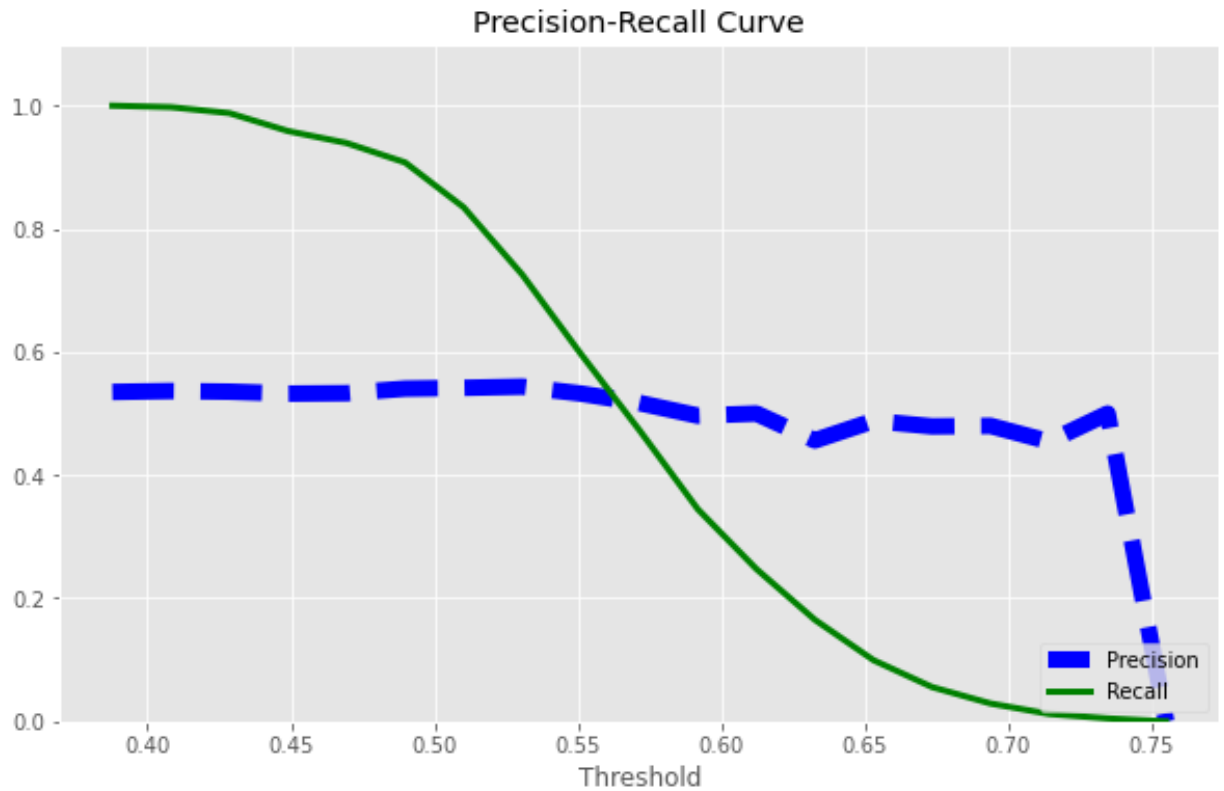
fig = plt.figure(figsize=(10, 6))
def plot_precision_recall_vs_threshold(precisions, recalls, thresholds):
    plt.plot(thresholds, precisions[:-1], "b--", linewidth=8, label="Precision")
    plt.plot(thresholds, recalls[:-1], "g-", linewidth=3, label="Recall")
    plt.xlabel("Threshold")
    plt.legend(loc="lower right")
    plt.title('Precision-Recall Curve')
    #plt.xlim([0, 1])
    plt.ylim([0, 1.1])

plot_precision_recall_vs_threshold(precisions, recalls, thresholds)
plt.show()

threshold_optimal = -1
for i in range(len(precisions)):
    if(precisions[i] == recalls[i]):
        threshold_optimal = thresholds[i]

print("Optimal Threshold: ", threshold_optimal)

```



Optimal Threshold: 0.7551020408163265

```
In [19]: print("Performance Measures Based on the Default Threshold:\n")

y_train_pred = cross_val_predict(knn, X_train, y_train, cv=5)

# Precision, Recall, F1 Score and Confusion Matrix for the Default Threshold 0.5
precision_train = precision_score(y_train, y_train_pred)
print("Precision (Default Threshold 0.5) = %f" % precision_train)

recall_train = recall_score(y_train, y_train_pred)
print("Recall (Default Threshold 0.5) = %f" % recall_train)

f1_train = f1_score(y_train, y_train_pred)
print("F1 Score (Default Threshold 0.5) = %f" % f1_train)

print("Confusion Matrix (Default Threshold 0.5)\n", confusion_matrix(y_train, y_train_pred))

print("\n-----\n")
print("Performance Measures Based on the Optimal Threshold (from Precision-Recall Curve):")
```

```

# Precision, Recall, F1 Score and Confusion Matrix for different thres
hold

t = threshold_optimal # optimal threshold from precision-recall curve

# Compute predictions based on new t by using the following method:
# - Get the probability of the positive class from the 2nd column[:,
1]
# - If that probability is greater than or equal to t, then the test
data belongs to the positive class
y_train_predicted_new = (cross_val_predict(knn, X_train, y_train, meth
od="predict_proba", cv=3)[:,-1] > t).astype(int)

precision = precision_score(y_train, y_train_predicted_new)
print("\nPrecision (Threshold %.2f) = %f" % (t, precision))

recall = recall_score(y_train, y_train_predicted_new)
print("Recall (Threshold %.2f) = %f" % (t, recall))

f1 = f1_score(y_train, y_train_predicted_new)
print("F1 Score = (Threshold %.2f) = %f" % (t, f1))

print("Confusion Matrix (Threshold %.2f)" % t)
print(confusion_matrix(y_train, y_train_predicted_new))

```

Performance Measures Based on the Default Threshold:

```

Precision (Default Threshold 0.5) = 0.552674
Recall (Default Threshold 0.5) = 0.827670
F1 Score (Default Threshold 0.5) = 0.662779
Confusion Matrix (Default Threshold 0.5)
[[ 84 276]
 [ 71 341]]

```

-----

Performance Measures Based on the Optimal Threshold (from Precision-Recall Curve):

```

Precision (Threshold 0.76) = 0.000000
Recall (Threshold 0.76) = 0.000000
F1 Score = (Threshold 0.76) = 0.000000
Confusion Matrix (Threshold 0.76)
[[360   0]
 [412   0]]

```

```
In [20]: #The accuracy of the model
test_accuracy_knn = knn.score(X_test, y_test)
print("\nTest Accuracy: ", test_accuracy_knn)

# No. of Correct Predictions
y_test_predicted = knn.predict(X_test)
print("\nNo. of correct predictions (Test): %d/%d" % (np.sum(y_test_predicted == y_test), len(y_test)))

# Confusion Matrix
print("\nConfusion Matrix (Test Data):\n", confusion_matrix(y_test, y_test_predicted))

precision = precision_score(y_test, y_test_predicted)
print("Test Precision = %f" % precision)

recall = recall_score(y_test, y_test_predicted)
print("Test Recall = %f" % recall)

f1 = f1_score(y_test, y_test_predicted)
print("Test F1 Score = %f" % f1)
```

Test Accuracy: 0.48704663212435234

No. of correct predictions (Test): 94/193

Confusion Matrix (Test Data):

[[18 79]

[20 76]]

Test Precision = 0.490323

Test Recall = 0.791667

Test F1 Score = 0.605578

```
In [21]: # Get the 2nd column of the matrix of predicted probabilities for each
data point
# The 2nd column stores the probabilities of the positive class
y_scores_test = cross_val_predict(knn, X_test, y_test, method="predict
_proba", cv=3)[: , 1]

fpr_test, tpr_test, thresholds_test = roc_curve(y_test, y_scores_test)

print("\nFPR FPR & TPR for Various Threshold Values:")
print("FPR: ", fpr_test)
print("TPR: ", tpr_test)
print("\nThresholds: ", thresholds_test)
```

FPR FPR & TPR for Various Threshold Values:


```
FPR: [0.          0.01030928 0.04123711 0.08247423 0.18556701 0.3092
7835
0.42268041 0.5257732  0.70103093 0.86597938 0.97938144 0.97938144
1.          1.          ]
TPR: [0.          0.          0.02083333 0.03125    0.0625     0.1562
5
0.35416667 0.53125    0.6875     0.82291667 0.94791667 0.96875
0.98958333 1.          ]
```

```
Thresholds: [1.6122449  0.6122449  0.59183673 0.57142857 0.55102041
0.53061224
0.51020408 0.48979592 0.46938776 0.44897959 0.42857143 0.40816327
0.3877551  0.34693878]
```



```
In [22]: plt.style.use('ggplot')

fig = plt.figure(figsize=(10, 6))
plot_roc_curve(fpr_test, tpr_test)
plt.show()
```



```
In [23]: # Area under ROC curve
roc_auc_score(y_test,y_scores_test)
```

Out[23]: 0.4463058419243986

```
In [24]: pd.crosstab(y_test, y_test_predicted, rownames=['True'], colnames=['Pr
edicted'], margins=True)
```

Out[24]:

Predicted	0	1	All
True			
0	18	79	97
1	20	76	96
All	38	155	193

```
In [25]: #The accuracy of the model
test_accuracy_knn = knn.score(X_test, y_test)
print("\nTest Accuracy: ", test_accuracy_knn)

# No. of Correct Predictions
y_test_predicted = knn.predict(X_test)
print("\nNo. of correct predictions (Test): %d/%d" % (np.sum(y_test_predicted == y_test), len(y_test)))

# Confusion Matrix
print("\nConfusion Matrix (Test Data):\n", confusion_matrix(y_test, y_test_predicted))

precision = precision_score(y_test, y_test_predicted)
print("Test Precision = %f" % precision)

recall = recall_score(y_test, y_test_predicted)
print("Test Recall = %f" % recall)

f1 = f1_score(y_test, y_test_predicted)
print("Test F1 Score = %f" % f1)
print(classification_report(y_test, y_test_predicted))
```

Test Accuracy: 0.48704663212435234

No. of correct predictions (Test): 94/193

Confusion Matrix (Test Data):

```
[[18 79]
```

```
[20 76]]
```

Test Precision = 0.490323

Test Recall = 0.791667

Test F1 Score = 0.605578

	precision	recall	f1-score	support
0	0.47	0.19	0.27	97
1	0.49	0.79	0.61	96
accuracy			0.49	193
macro avg	0.48	0.49	0.44	193
weighted avg	0.48	0.49	0.44	193

In [ ]: